

## Data

- Dataset: 20 Newsgroups (CMU subset).
- Split: 50% / 50% per class (stratified).
- Examples: Train = 9,998, Test = 9,999.
- Vocabulary (training): 105,428 unique tokens.
- Preprocessing: lowercase, regex alphanumeric tokenization, optional small stopword removal.

## Method

- Tokenizer: simple regex-based alphanumeric token extraction; stopwords may be removed in code.
- Model: Multinomial Naive Bayes implemented from first principles.
  - Class prior:  $P(y) = N_y / N$ .
  - Likelihood with Laplace smoothing:  
$$P(w | y) = (\text{count}(w, y) + \alpha) / (\text{total\_tokens}_y + \alpha \cdot V)$$
where  $\alpha = 1.0$  and  $V$  is the vocabulary size.
- Prediction: sum  $\log(P(y)) + \sum \log(P(w | y))$  over tokens to avoid underflow; predict argmax over classes.
- Implementation detail: token counts and log-probabilities are computed directly in Python.

The core implementation lives in the src/ directory

- `src/data/io.py`
  - I/O helpers for the dataset. Exposes `load_labeled_docs(root) → list of (label, text) pairs` and `discover_labels(root) → list of class folder names`. Keeps file reading separate from preprocessing/model logic.
- `src/text/processing.py`
  - Tokenizer & preprocessing. Main interface `preprocess(text, remove_stop=True)` (and `prepare_text`) performs: lowercase → regex alphanumeric token extraction → optional compact stopword removal → returns token list. This is the place to change tokenization / stopword behavior.
- `src/model/nb_patch.py` (or `src/model/nb.py`)
  - From-scratch Multinomial Naive Bayes implementation:  
`StudentMultinomialNB` with `train(docs, labels)` and `infer(docs)` methods.
    - Builds vocabulary (`vocab_`) from training data and per-class token counts.
    - Computes class log-priors and token log-probabilities with Laplace smoothing (`alpha`).
    - Predicts by summing log-priors + per-token log-likelihoods and selecting the highest score.
    - Exposes scikit-like attributes for reporting: `vocab_`, `classes_`, `class_log_prior_`, etc.
    - Uses `collections.Counter` and plain Python arithmetic no ML libraries.
- `src/eval/metrics.py`
  - Evaluation utilities: `accuracy`, `confusion`, `per_class_prf`, and `tsv_confusion`. These produce the overall metrics, confusion matrix (TSV), and per-class precision/recall/F1 used in the report

## Experimental details

- Split reproducibility: seed = 42.
- Training aggregates per-class token counts, computes log-likelihoods, and caches unseen-token log-probabilities.
- Evaluation metrics: Accuracy, Macro F1, Micro F1, per-class precision/recall/F1, and confusion matrix.

## Results (main figures)

- Accuracy: 0.8702
- Macro F1: 0.8661
- Micro F1: 0.8702
- Train examples: 9,998 — Test examples: 9,999
- Vocabulary size: 105,428

## Sample per-class performance (selected examples)

- rec.sport.hockey — F1  $\approx$  0.967
- sci.crypt — F1  $\approx$  0.960
- comp.graphics — F1  $\approx$  0.794
- talk.religion.misc — F1  $\approx$  0.660

(Full per-class metrics and confusion matrix are available in the project reports/directory.)

## 7. Observations & short error analysis

- The classifier does particularly well on narrowly-scoped topics (sports, some science categories) where topic-specific vocabulary strongly signals the class.

- Errors concentrate among semantically close groups (for example, several comp. classes and some political/religion subcategories). Overlap in technical terms and common vocabulary across these topics increases misclassification rates.
- Long-tailed tokens and rare proper nouns can skew likelihoods for individual classes. Possible improvements (not required for this assignment) include vocabulary pruning, stemming, or TF-IDF weighting prior to classification.

## Reproducibility & how to run

Create the split :

run the tolerant splitter:  
python auto\_split.py

Train and evaluate:

python main.py --data\_root data --alpha 1.0 --report reports/report.md

There is a convenience script to run both steps:

./run\_all.sh